



Per-step latency percentiles

NVIDIA L4 distribution (n=479 steps; x-axis zoomed ≥ 30 ms)

L4: real validation step streams; M4 Pro: real holdout step streams; warm, single-stream, incremental KV cache, clocks not locked.
Header prefill (~ 40 ms on L4) amortizes to $\sim +2.9$ ms/step at the 14-step median holdout trajectory (median from results/v2_holdout_step1100.json).
StepShield's native timing metric is EIR (steps); ms figures are from our own harness.