


Leakage-free threshold calibration: every decision made by a threshold that never saw that trajectory or its twin



DEPLOYED θ = max of ALL 108 clean scores = 0.946597 $\rightarrow$ forward certificate $E[\text{FPR}] \leq 1/109 = 0.917\%$ (exchangeability); an $\text{FPR} \leq 0.5\%$ certificate is impossible at $n_c=108$ (needs ≥ 199).
Leakage-free recall 90.7% (θ never saw a rogue score; same 108 rogues as the CV table — correlated, not independent). In-sample FPR is 0 by construction — NOT evidence.